

# Chapter 1

## Implementation and Evaluation

### 1.1 Evaluation of Other HWR Systems

The recognition rates reported in the literature are shown in figure 1.1 borrowed from ? (?). As a general trend it can be noted that the recognition rate of most systems lies between 85% and 95%. They believe that it is possible to achieve a recognition rate of up to 98% for regular scripts. On fluent scripts, however, they regard it as difficult to achieve a recognition rate above 90%. The systems marked with an asterisk in figure 1.1 perform recognition for Chinese or Japanese characters. All of their recognition rates lie below 90%. That performance measure sets a general context in which the prototype system developed in this work might be arranged. However, the results of other systems are only comparable to the overall recognition accuracy of the present system. Due to the different technical background and the aim of the system it is difficult to draw conclusions from the sheer comparisons of numbers. (See section 1.3.2.2 for more information.)

### 1.2 E-Learning Module Evaluation

Generally, there are two main directions in the evaluation methodology for e-learning applications: the educationalist's approach and the software developer's approach. Therefore, the evaluation of an e-learning system is a complex task and requires optimisation work on the account of both the conceptual designer of an e-learning application as well as the software developer.

The e-learning part of the prototype is a sample module that is used to exemplify one usage scenario of the HWR engine. It mainly shows plausibility of the approach. A detailed analysis of the e-learning application using ISO9126 (?) is not useful, since the e-learning part of the software has not been optimised in any way. The focus of the thesis is not to implement an e-learning application, but to create an analytical handwriting recognition engine. It might be a prospect for future work<sup>1</sup> to optimise the e-learning part and build a fully-developed e-learning application for Japanese characters, but that would be outside the scope of this thesis. For these reasons, there will not be an evaluation of the e-learning module.

### 1.3 Evaluation of the HWR Engine

#### 1.3.1 General Considerations for Evaluation

The performance of a recognition system can be measured in terms of speed, accuracy and memory requirement. While statistical systems offer high speed but have large requirements on memory, structural methods have lower speed but require only a smaller memory (?). The system developed and evaluated in this thesis is a structural system. It can be expected that the system has a relatively low performance speed, but moderate memory requirements. Additionally, since the system is not just a structural but an analytical recognition system, the speed might be even lower.

The factors memory requirements and speed are not of great interest in the context of this system. The system is an on-line system, but it exclusively performs single character recognition. The focus lies a lot more on a detailed analysis of one character, rather than the high-speed recognition of a stream of characters. The system is an interactive system. The recognition of a single character is an in-depth analysis of the structure of that character and returns a profound feedback to the user. Especially in an interactive learning context, the user is supposed to work with the system's feedback. Recognition speed is interesting for evaluation if the user enters a stream of characters. In that case recognition speed can be expressed as a factor that relates recognition

---

<sup>1</sup>See section ??

Source	Method	#category	Style	#learning	#test	Rec. rate
Liu'91 [75]	struct	6,763	flu-regular	N/A	N/A	90%
Kawamura'92 [45]	statis	2,965	careful/free	380 PC	20 PC	94.51/91.78%
Lin'93 [65]	struct	5,400	regular	1 PC	10 PC	87.4%
Liu'93 [76]	struct	13,000	flu-regular	N/A	N/A	93%
Hamanaka'93 [28]	statis	1,064	regular	54,028	52,944	95.1%
Chou'94 [16]	struct	5,401	regular	3 PC	17 PC	94.88%
Wakahara'95 [133]	struct	2,980	careful/free	120 PC	36 PC	97.6/94.1%
Lay'96 [60]	struct	5,401	regular	N/A	5 PC	96.35%
Kim'96 [47]	struct	1,800	free	4 PC	6 PC	93.13%
*Nakagawa'96 [88]	struct	3,345	fluent	N/A	11951 PW×30	80–90%
Chou'96 [18]	struct	5,401	flu-regular	5 PC	15 PC	93.4%
Wakahara'97 [134]	struct	2,980	careful/free	120 PC	36 PC	98.4/96.0%
Kim'97 [48]	HMM	1,800	free	4 PC	6 PC	90.3%
Zheng'97 [149]	FARG	3,755	regular	N/A	6 PC	98.8%
Xiao'97 [139]	struct	3,755	flu-regular	35 PC	3 PC	93.9
Nambu'98 [94]	struct	3,942	flu-regular	200 PC	200 PC	89.7%
Kuroda'99 [58]	statis	1,000	regular	25 PC	10 PC	94.34%
*Okamoto'99 [99]	statis	3,345	fluent	11951 PW×40	11951 PW×41	86.32%
*Yasuda'99 [143]	HMM	3,057	fluent	10038 PW×100	10038 PW×20	85.89%
*Tanaka'99 [119]	combined	3,356	fluent	Nakayosi	Kuchibue	87.6%
Zheng'99 [150]	sta-struct	3,755	mixed	100 PC	7 PC	95.52%
*Akiyama'00 [1]	struct	3,345	fluent	6690 PW×150	11951 PW×3	88.58%
Shin'02 [113]	struct	2,965	regular	90 PC	24 PC	99.28%
*Kitadai'02 [53]	struct	3,345	fluent	9309 PW×163	11951 PW×120	87.2%
Tokuno'02 [121]	HMM	1,016	fluent	50,986	42,718	92.0%
Velek'02 [129]	combined	3,036	fluent	3,669,089	54,927	94.14%
Nakai'02 [92]	HMM	1,016	fluent	34 PC	34 PC	93.1%
Matic'02 [79]	neural	4,400	N/A	80 PC	20 PC	97.3%
Rowley'02 [106]	struct	6,847	natural	5 million	85,655	94.45%

PC: per category

PW: per writer

\*Tested on TUAT Kuchibue database

Figure 1.1: Recognition rates reported in the literature

speed with input speed. ? (?) state that on-line recognition systems need only be fast enough to keep up with the writing. Further, they report average writing rates of 1.5-2.5 characters/s for English alphanumerics and 0.2-2.5 characters/s for Chinese characters. In a system that performs single character recognition the user has the impression of instant recognition.

Memory requirements are negligible in this system for a similar reason. The recognition of one character does not require much memory compared to the recognition of a stream of characters. Additionally, the main system engine does not run on a small mobile computing device with low memory capacity, but runs as a service on a standard PC. With the advanced memory capacities of today memory is not an issue. Nevertheless, even mobile devices are equipped with enough memory to enable the system to perform analytical character recognition.

For the reasons given above, the evaluation of the analytical handwriting recognition engine will be limited to different types of accuracy measurements.

### 1.3.2 Development of Appropriate Evaluation Metrics

#### 1.3.2.1 Choice of Evaluation Subjects

It is difficult to perform an accuracy evaluation of a recognition system that can be compared to other systems. The methods the systems use in order to perform their recognition are diverse. There is always a trade-off between robustness, performance and accuracy. The prototype that is subject to this evaluation performs a different task than most of the other recognition systems. It analyses the characters not only for the purpose of recognition, but attempts to create feedback on how well the input matched the character model in structural terms. That means concretely, the system can analyse an input with an expected result and can perform the




Figure 1.2: Similar characters that can be confusing to learners: Five different Kanji share the same Radical in the *tsukuri* position (on the right).

same for an unknown input by assuming the best match as the expected result. The analysis yields an output that includes more information than pure pattern matching. It rather includes structural linguistic information.

For example, the characters 堉, 𦵏, 𦵐, 𦵑 and 鉈 all share the substructure on the right: 花. Figure 1.2 shows the substructure on top and below five characters that are using the structure. The system analyses and distinguishes substructures. If the system is set up to recognise an input with an expected result the output contains a confidence value about the input quality concerning that character. Additionally, the error recognition module returns information about the substructures that were found in the input. The similarity of the characters from figure 1.2 is known to the recognition system because of the identical substructures. If the system identifies the input as a character with identical parts with respect to the expected character the output will contain the information as a substructure confusion error type.

This detailed analysis creates a unique requirement for evaluation. Not only recognition accuracy must be measured and compared to other systems, but a new evaluation for the recognition of substructures is needed. It seems reasonable to measure the recognition accuracy in a plain percentage of correctly recognised characters. Precision and recall are the correct measurement for the error recognition, where characters with incorrect substructures should be analysed. That way, true positives, false positives, true negatives and false negatives can be distinguished. Since the system can recognise pure substructures as well, it would be interesting to see the accuracy of a Radical recognition.

The lexicon is a sample lexicon that contains 50 characters. That circumstance has to be kept in mind when comparing the evaluation results to other systems. Creating a lexicon entry for a single character is a laborious task. It would be outside the scope of the thesis to create a large lexicon. Nevertheless, the characters have been chosen by graphic and semantic criteria in order to create more possibilities for confusion. There are many characters using the same key Radical, other characters that simply look similar, which might evoke difficulties for both the learner and the system. Three experiments will be conducted in order to evaluate the analytical handwriting recognition engine:

1. Overall recognition accuracy

## 2. Substructure recognition accuracy

## 3. Error recognition accuracy

The metrics details for these experiments will be presented in the following sections.

### 1.3.2.2 Overall Recognition Accuracy

The overall recognition accuracy (experiment 1) will be calculated as a measure of the  $n$ -best matches for a character input.  $N$  is defined as the lowest position in the list of  $n$ -best matches that will still be considered as a *match* of a character. That is, if  $N$  is set to 1, then only the best match will be considered. If  $N = 3$ , the best three matches will be taken into account. The overall recognition accuracy  $A_N$  is defined here as the weighted percentage of correctly recognised characters when taking the  $N$  best matches into account. The number of samples will be called  $S$ .  $r$  is an integer that marks the position in the list of  $n$ -best matches. The values for  $r$  that will be considered for the evaluation lie on the open interval  $]1, N[$ . If an input produces a confidence value for a character recognition high enough to be in a list position  $r \leq N$ , then  $\frac{1}{r}$  will be added to the number of correctly recognised characters. That means, if a character is the best match in the list of  $n$ -best matches, then it is ranked on top, therefore  $r = 1$ . Hence,  $\frac{1}{r} = \frac{1}{1} = 1$  will be added to the number of correct matches. If the correct character is the second best match,  $\frac{1}{2}$  will be added and so forth. If the input has been correctly recognised as the  $r$ -best match, but  $r > N$  or, alternatively, the input has not been recognised at all, nothing will be added.  $f$  is a helper function that returns a value according to the relation between  $r$  and  $N$ . The variable  $r_i$  denotes the position of the  $i^{\text{th}}$  input character in the list of  $n$ -best matches for that input sequence and character.  $r_i$  serves as an input value for  $f$ .  $A_N$  will be calculated as follows:

$$f(r, N) := \begin{cases} \frac{1}{r} & \text{if } r \leq N \\ 0 & \text{if } r > N \end{cases}$$

$$A_N := \frac{1}{S} \cdot \sum_{i=1}^S f(r_i, N)$$

An example for the application of that term would be to set  $N = 1$ . Only the best match will be considered for an input. If 100 input sequences are tested, then  $S = 100$ . If, for example, one of the  $S_i$  input sequences yields a high confidence value for the character the sequence is supposed to represent and the character becomes the most salient in the list of best matches and occupies the first position. Then  $r_i = 1$  and  $f(r_i, N) = f(1, 1)$  yields 1. Therefore, in the summation of all the  $f(r_i, N)$  the value 1 will be added for each correctly recognised character. Say, 69 of the 100 input characters were correctly recognised, then  $A_N = A_1 = \frac{1}{100} \cdot 69$ . That would equal a result of 69%. The key to this evaluation method is of course the value of  $N$ . Multiple experiments with different  $N$ -values will be conducted for best comparability.

**Baseline.** Different possibilities exist for the definition of a baseline for that experiment. A baseline based on pure randomness would be  $\frac{1}{C}$  with  $C$  representing the number of characters in the database. For  $C = 50$  the baseline would then be  $\frac{1}{C} = \frac{1}{50} = 0.02$ . That is already a very low baseline. The larger the lexicon grows the lower the baseline for evaluation. For a lexicon that contains around 2,000 characters and covers the Jōyō Kanji the baseline would come down to  $\frac{1}{C} = \frac{1}{2000} = 0.0005$ . This baseline is certainly not sufficient to measure the quality of a recognition system. However, a significant baseline can be defined with the formula for the computation of  $A_N$  and generous assumptions about the accuracy of the recognition:

Assume,  $N = 3$  and all the sample characters are among the first three matches in the list of matches, equally partitioned. Then one third would yield  $r = 1$  with  $f(r, N) = f(1, 3) = 1$ , another third yield  $r = 2$  with  $f(r, N) = f(2, 3) = \frac{1}{2}$ , and the last third yield  $r = 3$  with  $f(r, N) = f(3, 3) = \frac{1}{3}$ . That leads to a baseline ( $B$ ) calculation as follows:

$$\begin{aligned} B &:= \frac{1}{S} \cdot \left( \sum_{i=1}^{\frac{S}{3}} f(1, 3) + \sum_{i=1}^{\frac{S}{3}} f(2, 3) + \sum_{i=1}^{\frac{S}{3}} f(3, 3) \right) \\ &= \frac{1}{S} \cdot \left( \sum_{i=1}^{\frac{S}{3}} 1 + \sum_{i=1}^{\frac{S}{3}} \frac{1}{2} + \sum_{i=1}^{\frac{S}{3}} \frac{1}{3} \right) \\ &= \frac{1}{S} \left( \frac{S}{3} + \frac{S}{3 \cdot 2} + \frac{S}{3 \cdot 3} \right) = \frac{1}{3} + \frac{1}{6} + \frac{1}{9} = \frac{11}{18} = 0.61 \end{aligned}$$

The baseline value of 0.61 will be used for the overall recognition accuracy experiment.

### 1.3.2.3 Substructure Recognition Accuracy

The substructure recognition accuracy (experiment 2) will be measured in a similar way as the overall recognition accuracy. In fact the same metrics will be used. See the previous section for details on the metrics. The accuracy  $A_N$  will be measured as a percentage value of correct recognition results, where the recognition as second-best or third-best (generally  $n$ -best, with  $n > 1$ ) will be devaluated. This reduction of value is expressed mathematically by using the multiplicative inverse  $r^{-1}$  of the rank  $r$  in the list of  $n$ -best matches as a summand for the weighted sum of correct recognitions. The equations for calculating the accuracy have been given in section 1.3.2.2 on page 4.

The main difference is that many substructures occur in more than one character. In order to account for a substructure recognition will be regarded as *correct*, if the input yields the intended substructure from one of the characters that contain the substructure. That fairly generous interpretation of what is correct considers the fact that the substructures should be recognised for any of the characters. During the recognition process it does not matter, what character needs to be recognised the substructure recognition should work independently. In a future version of the database the substructures should be stored only once and then be shared among all characters that contain it. The *baseline* for the substructure recognition will be 0.61, the same number and using the same line of argument from section 1.3.2.2.

### 1.3.2.4 Error Recognition Accuracy

The error recognition accuracy (experiment 3) is more difficult to evaluate than the pure accuracy. The question arises, what kind of errors are expected to be found and what kind of error corrections should be provided. Error recognition here means that the system compares a known character with a given input stroke sequence. The input can contain flaws that do not represent the character appropriately. The prototype is expected to find the flaw and propose a correction. The evaluation method will be a classical *precision and recall* method. There are two possible cases in the input stroke sequence:

- There is an error in the input with respect to the desired character.
- There is no error in the input.

There are two possibilities for the system to react:

- The system identifies an error in the input with respect to the desired character and gives feedback on that error.
- The system does not identify an error in the input with respect to the desired character.

Those possibilities yield a  $2 \times 2$ -matrix: Table 1.1 gives an overview about the possible results combined with the reality of the sample.

	Error in input	No error in input
Found error (TP+FP)	true positive (TP)	false positive (FP)
No error found (FN+TN)	false negative (FN)	true negative (TN)

Table 1.1: Possible results as a basis for precision and recall analysis.

The precision  $P$  and recall  $R$  of the error recognition are computed with the standard formulas:

$$P = tp \cdot (tp + fp)^{-1}$$

$$R = tp \cdot (tp + fn)^{-1}$$

Additionally, the *F-measure* will be used, however only for an equally important interpretation of precision and recall. Therefore,  $\beta = 1$  will be fixed:

$$F_\beta = (1 + \beta^2) \cdot \frac{P \cdot R}{\beta^2 \cdot P + R}$$

$$F_1 = 2 \cdot \frac{P \cdot R}{P + R}$$

## 1.3.3 Experimental Results

Having presented the experimental metrics in the previous sections, the experimental results and the conclusions drawn will be presented in this section.

### 1.3.3.1 Overall Recognition Accuracy Results

For the experiment concerning the overall recognition accuracy two test writers wrote all 50 characters from the database as an input sequence for the system. That process yielded 100 input stroke sequences. The input was stored in order for the experiment to be repeated without physically re-entering the characters. Five sample characters from each writer haven been removed from the test set for detailed analysis. With five characters from each writer in the development set, there were 90 characters for the test set. The test set characters have been used only for testing, not for detailed analyses. The development set characters have been tested as well and the recognition process as well as the recognition result have been examined in order to improve the recognition engine.

The final evaluative run of the test set has been repeated 10 times, with different values for  $N$ . All of the integers from 1 to 10 have been used. The weighted number of recognised characters (WNCR) changes according to the number of characters recognised in each rank in the list of  $n$ -best matches. Table 1.2 shows the results in greater detail. With each incrementation of  $N$ , that is with decreasing strictness the recognition results rise. The range of the weighted accuracy results goes from 70% for  $N = 1$  up to 81% for  $N = 5$ . Taking into account even greater ranks up to  $N = 10$  in the list of  $n$ -best matches increased the result slightly to 82%. This development can also be seen in the chart in figure 1.3.

N	$A_N$	WNRC/90	1	2	3	4	5	6	7	8	9	10
1	0.70	63.00	63	-	-	-	-	-	-	-	-	-
2	0.74	67.00	63	8	-	-	-	-	-	-	-	-
3	0.77	69.33	63	8	7	-	-	-	-	-	-	-
4	0.80	71.83	63	8	7	10	-	-	-	-	-	-
5	0.81	72.63	63	8	7	10	4	-	-	-	-	-
6	0.81	72.80	63	8	7	10	4	1	-	-	-	-
7	0.81	73.23	63	8	7	10	4	1	3	-	-	-
8	0.82	73.35	63	8	7	10	4	1	3	1	-	-
9	0.82	73.58	63	8	7	10	4	1	3	1	2	-
10	0.82	73.68	63	8	7	10	4	1	3	1	2	1

Table 1.2: Overall recognition results.

**Result Interpretation.** As mentioned earlier, the characters bear a lot of potential for confusion, since many of the database characters have been chosen by their key Radical or by similar graphical appearance to other characters already in the database. That special choice of characters seems to explain why there is such an increase in accuracy when more ranks are taken into account. The similar characters will be confused by the system and appear in the lower ranks (higher  $r$ -values) in the list of  $n$ -best matches. Even though all the results, even the one with  $N = 1$  are sizably above the baseline of 61%. But generally, the accuracy results are considerably lower than those of other systems (see section 1.1). However, it has to be noted that the aim of this system is not to provide a handwriting recognition highly optimised for speed and accuracy. Hence, there is little comparability for these results. The prototype of the analytical handwriting recognition fulfils an additional task that the other systems do not aim at. That obviously leads to a trade-off concerning pure recognition performance. The aim is to create a deep analysis of a character structure and be able to provide feedback on the input. These abilities of the system will be evaluated in the following sections.

The analysis of recognition errors among the development set characters revealed that there are two main types of errors:

1. **Different stroke order**
2. **Similar characters**
3. **Segmentation ambiguity**

The type 1 and type 2 errors probably occur for all handwriting systems, but especially for overlaid or continuous handwriting. Those types of error might be resolved if the recognition engine was using a statistical model like a *Hidden Markov Model*. The segmentation ambiguity is a problem that occurs in overlaid handwriting recognition for full characters. In analytical handwriting recognition the problem occurs for substructures. The beginning and the end of a substructure can only be found by completely analysing the substructure. That is a time consuming process but must be solved.




Figure 1.3: Overall recognition results chart

### 1.3.3.2 Substructure Recognition Accuracy Results

For the evaluation of the recognition of substructures it would have been desirable to use the same input data with which the overall evaluation had been conducted. That was not possible, because it would have required to analyse the input data and split it into substructures manually. The input data for the evaluation of the substructures has been provided by three writers each of whom entered 25 substructures, i.e. a total of 75 Radicals and Graphemes. That means 75 different tokens, not 75 different types. But each type has only been written once by each writer. This input set is a pure test set it has not been split or used for development work. The results of the substructure recognition evaluation might be more comparable to the results of other system. The metrics for the substructure recognition are essentially the same as for the overall evaluation. The difference is that only the first three entries in the  $n$ -best list will be considered. The structures are smaller, need less analysis, do not have further substructures except the individual strokes. As expected, the evaluation results are higher than those of the overall recognition accuracy. The general trend can be seen in figure 1.4, the complete data is laid out in table 1.3.

N	$A_N$	WNRC/75	1	2	3
1	0.85	64	64	-	-
2	0.89	67	64	6	-
3	0.90	67.33	64	6	1

Table 1.3: Substructure recognition results.

**Result Interpretation.** As expected, the accuracy of the recognition results for the substructures was higher than the overall character recognition accuracy. The substructure recognition is more robust as it does not have a segmentation problem. Each input sequence is known to be a full substructure. During the recognition process the database needs to be searched for substructures. If a matching structure is found the process is completed. The same pattern like in overall recognition accuracy can be observed concerning the lower positions in the list of  $n$ -best matches.

### 1.3.3.3 Error Recognition Accuracy Results

The error recognition evaluation required more manual work than the overall recognition or substructure recognition evaluation. This is due to the fact that errors had to be produced intelligently in order for the error recognition to pick them up. Therefore, this evaluation type has to be seen more as a qualitative analysis, rather than a pure quantitative evaluation.

**Test Data Creation.** One writer created 30 characters, including 20 characters with errors, i.e. 20 characters where one Radical had been replaced with another one or with garbled input. Figure 1.5 shows the characters 没 (Jap. pron. *botsu*; Eng. 'to sink') and 海 (Jap. pron. *umi*; Eng. 'sea') both in print and handwritten versions. The two characters share the same key Radical. It is the 'water'-Radical on the left side (marked in red). On the right side, the lower Radical is fairly distinct and therefore no cause of confusion. In fact 海 (*umi*/'sea') contains the 'mother'-Radical 母 (marked in purple), while 没 (*botsu*/'to sink') contains the 'again'-Radical 又 (marked in purple), which looks very similar to the 'father'-Radical 父 (not contained in the figure). This mother/father-analogy has nothing to do with the true meaning of the characters, it can just serve as a mnemonic for a learner. This fact could be a cause for semantic confusion due to the mnemonic trick, but clearly those two Radicals offer little features for a graphical confusion between the two. The upper Radical on the right in both characters is a small two-stroke Radical that can easily be confused graphically with another small two-stroke Radical (marked in green). This danger of confusion exists both for a human and for the character recogniser. Non-existent blended characters like the one shown in figure 1.6 have been created as erroneous input for the error recognition evaluation test set. In case of a confusion of the two Radicals in the upper right corner of the character the user might produce an input like the one in figure 1.6.

**Error Recognition Result Data Set.** Table 1.4 shows the results of the error recognition test run. In section 1.3.2.4 it has been argued that the correct evaluation method for the error recognition was a precision and recall measurement. The problem with a pure precision and recall analysis is that the results in table 1.4 are not 2-way but 3-way. The unrecognised characters could be regarded as incorrect or correct. If there was an error in the input character and the character could not be recognised, this could be regarded as a *true positive*. This would increase the error recognition performance undeservedly. On the other hand, if the input was a correct character, then assuming the result to be an error message would add to the number of *false*




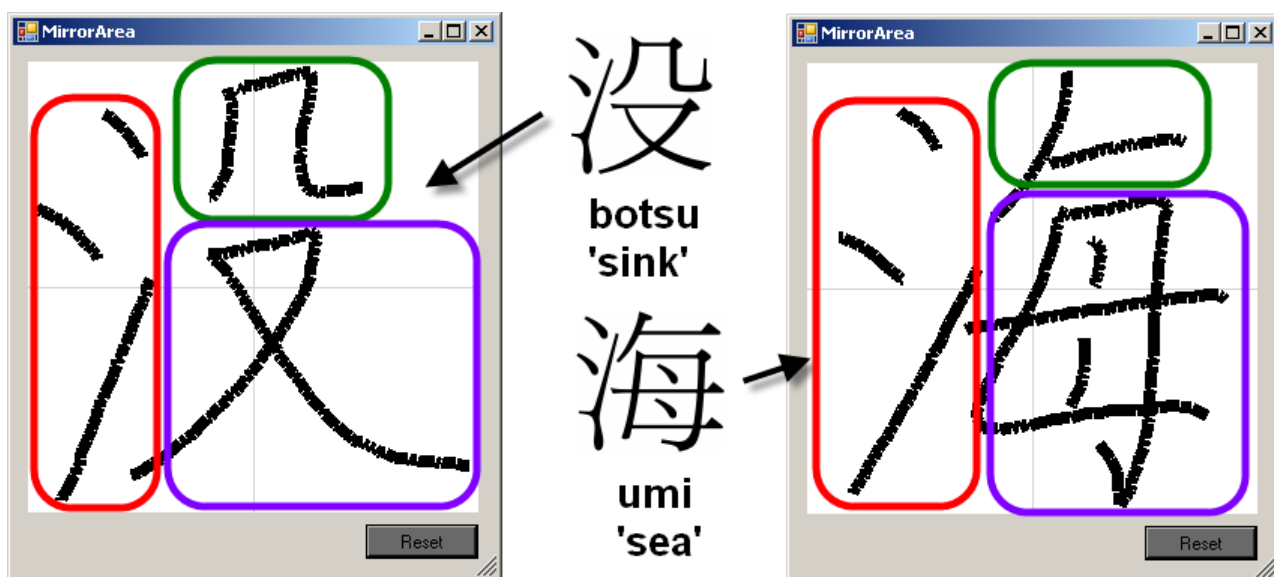


Figure 1.4: Substructure recognition results chart

Figure 1.5: Characters 没 (*botsu*, 'sink') and 海 (*umi*, 'sea')

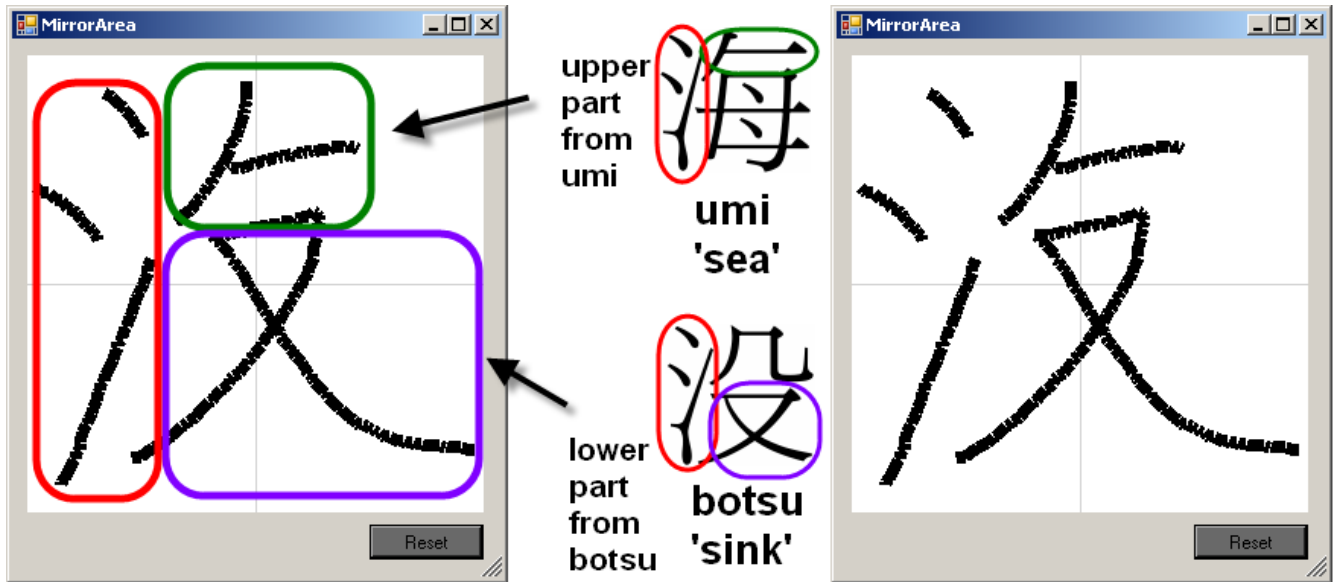


Figure 1.6: Confusion between the characters 没 (*botsu*, 'sink') and 海 (*umi*, 'sea')

*positives*, lowering the performance result of the error recognition. The error recognition is a sub-component of the character recognition engine. It cannot work independently from it. Therefore it would not be a fair judgement to include the unrecognised characters at all into the evaluation of the error recognition. Firstly, in order to ensure a fair evaluation but also in order to provide a result set qualified for a 2-way evaluation, the unrecognised characters are exempt from the evaluation of the error recognition. Table 1.5 shows the data set that was used for the calculation of precision and recall without the unrecognised characters. The total input set consisted of 30 characters, 20 with errors in the characters, 10 correct characters. After exempting the unrecognised characters from the evaluation there is a total of 25 characters, 17 of which containing errors, 8 correct characters.

Total of 30 chars (20/10)	Error in input	No error in input
16 found error (TP+FP)	14 (TP)	2 (FP)
9 no error found (FN+TN)	6 (FN)	3 (TN)
5 unrecognised (?)	3 (?)	2 (?)

Table 1.4: Error recognition results as a basis for precision and recall analysis.

Total of 25 chars (17/8)	Error in input	No error in input
16 found error (TP+FP)	14 (TP)	2 (FP)
9 no error found (FN+TN)	6 (FN)	3 (TN)

Table 1.5: Error recognition results excluding the unrecognised characters as a basis for a 2-way precision and recall analysis.

**Document created on Tuesday 22<sup>nd</sup>  
June, 2010 at 10:56**